

# NclaveOS Agentic Operating System

*An Early-Access Open Source Project Made by Exxeta*

| *KI-Agenten für Teams — governed, auditierbar, self-hosted*



OPA Policy Gate



Multi-User RBAC



Secret Isolation



Approval Gate



Audit Trail



Scheduled Tasks

## Das Problem

---

### Jeder baut sich eigenen Agenten Stack

Jeder konfiguriert lokal — keine gemeinsamen Skills, keine einheitlichen Policies.

### Kein Audit

Was hat der Agent auf Prod gemacht? Weiß keiner.

### Secrets im Prompt

API-Keys landen im LLM-Kontext, in Logs, in der Bash-History.

### Kein Gate

Der Agent kann Prod-Ressourcen löschen — nichts hält ihn auf.

# Introducing: NclaveOS Agentic Operating System

---

*From personal automation to governed agent infrastructure*

- **zentral gehostetes** Agenten-Management
- Agenten planen und führen Commands aus
- Skills und Systemzugriffe zentral verwaltet
- deterministische Guardrails zentral konfigurierbar
- opt. Approval-Gate vor Commands
- Audit Trails
- Cron-Jobs
- [...]

## Lokaler Assistent vs. governed Platform

---

| Fähigkeit                   | Claude Desktop / MCP | nclaveOS   |
|-----------------------------|----------------------|------------|
| Multi-User mit Rollen       | ×                    | ✓          |
| Policy-gated Execution      | ×                    | ✓ OPA Rego |
| Secret-Isolation            | ×                    | ✓          |
| Human Approval Gate         | ×                    | ✓          |
| Scheduled Tasks             | ×                    | ✓ Cron     |
| Audit Trail                 | ×                    | ✓          |
| Gemeinsame Skill-Bibliothek | ×                    | ✓ Git      |
| Self-hosted                 | manchmal             | ✓ immer    |

## Wie es funktioniert

---

|                                                                                          |   |                                                                                                 |   |                                                                                              |                                                                                       |                                                                                              |   |                                                                                                 |
|------------------------------------------------------------------------------------------|---|-------------------------------------------------------------------------------------------------|---|----------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------|---|-------------------------------------------------------------------------------------------------|
|  Prompt | → |  Planner (LLM) | → |  OPA Policy |  → |  Executor | → |  Audit Trail |
|                                                                                          |   | ↑ result                                                                                        |   |  deny      |                                                                                       |                                                                                              |   |                                                                                                 |

| *Planner → OPA → Executor → result zurück an Planner — Loop bis done oder failed*

## Skills

---

```
name: kubectl-readonly
description: |
  Read-only k8s CLI.
  Only get, describe, logs, top.
policy: |
  allowed := {"get","describe","logs","top"}
  allow {
    input.argv[0] == "kubectl"
    allowed[input.argv[1]]
  }
env:
  - KUBECONFIG_TOKEN
```

### Admins definieren:

- welche Tools verfügbar sind
- welche Commands erlaubt sind
- welche Secrets injiziert werden

### Verteilt via Git-Repository →

ein Sync-Klick, alle Teams haben dieselben Skills.

## OPA als Ausführungsgate

---

| Command                               | OPA                           | Ergebnis                                                                                     |
|---------------------------------------|-------------------------------|----------------------------------------------------------------------------------------------|
| <code>kubectl delete pod x</code>     | <code>delete ∉ allowed</code> |  blocked  |
| <code>kubectl get pods -n prod</code> | <code>get ∈ allowed</code>    |  executed |
| <code>rm -rf /</code>                 | kein Skill matched            |  blocked  |

Globale Policy + optionale **per-Skill-Policy** — in-process, kein Netzwerk-Round-Trip.

## Secret-Isolation

---

|                                                                                               | sieht                                                         |
|-----------------------------------------------------------------------------------------------|---------------------------------------------------------------|
|  LLM         | <code>\${GITHUB_TOKEN}</code> — Placeholder                   |
|  Audit Trail | <code>\${GITHUB_TOKEN}</code> — Placeholder                   |
|  Subprocess  | <code>GITHUB_TOKEN=ghp_xxxx</code> — echter Wert              |
| <code>secrets.json</code>                                                                     | <code>chmod 600</code> , gitignored — nur der Server liest es |

Secrets verlassen **nie** den Subprocess — nicht im LLM-Kontext, nicht in Logs, nicht in der Run-History.



# Browser-UI

---

## Nutzer

-  Chat-Interface
-  Approve / Deny per Command
-  Run-History durchsuchen
-  Skills pro Chat aktivieren
-  Scheduled Tasks (Cron)

## Admins

-  User-Verwaltung + Rollen
-  LLM-Endpoint & Modell
-  Skills (inkl. Guardrails) erstellen, testen und ausrollen
-  Approval global ein/ausschalten

## Deployment

---

```
docker compose up -d
```

|                    | Option A              | Option B                             |
|--------------------|-----------------------|--------------------------------------|
| <b>LLM-Backend</b> | OpenAI / Azure OpenAI | Ollama (lokal)                       |
| <b>Persistenz</b>  | JSON Files (default)  | MongoDB via <code>MONGODB_URI</code> |
| <b>Port</b>        | <code>:8081</code>    | <code>:8081</code>                   |

## Roadmap

---

| Feature                  | Status                                                                                       |
|--------------------------|----------------------------------------------------------------------------------------------|
| Per-Chat Model Selection |  In Planung |
| Webhook-Trigger          |  In Planung |
| Kubernetes Helm Chart    |  In Planung |
| Skill Marketplace        |  Idee       |

# Mitmachen

---

## Nutzer

1. `docker compose up -d`
2. Skill laden
3. Run starten

## Contributor

1. Fork + Feature Branch
2. `pytest` / `npm run test`
3. Pull Request

*MIT License · © 2026 Exxeta AG · [github.com/exxeta/nclaveos](https://github.com/exxeta/nclaveos)*